

# FYS-STK 3155 project 2

Theodor Midtbø Alstad

14th November 2020

## Abstract

This report compares a neural network to numerical stochastic gradient descent linear regression, analytical Ridge linear regression, and logistic regression for classification. It finds that the neural network outperforms all compared methods for linear regression, and is outperformed by the logistic regression for classification because it had more features implemented.

## 1 Introduction

A neural network is a very flexible data analysis tool that can grant a large amount of knowledge about correlations between features when applied to a large enough volume of data, but it requires a lot of optimization and parametrization. It is therefore important to ascertain that it does perform well enough to be worth the time and resources building and training it takes.

In this project, I first implement and study Stochastic Gradient Descent where I compare it as a simple numerical approach with analytical approaches in OLS and Ridge and use them as a basis to determine how well the neural network performs linear regression. I implement and study a logistic regressor as a basis to determine how well the neural network performs classification. Finally, I implement and study a Neural Network as the most detailedly implemented algorithm, and studied determined whether it outperformed the previously mentioned algorithms for linear regression and classification.

This report contains a regression analysis on the Franke function (subsection 1.2) and a classification analysis on the MNIST dataset (subsection 1.3) as a basis for evaluating the Neural Network.

The scripts for the Stochastic Gradient Descent and Linear Regression are not optimized for user friendliness, as they were created to form a basis for comparison for the neural network.

### 1.1 The structure of the source code and filesystem

The source code is split into 3 main scripts and 3 side scripts.

Main scripts:

- SGD.py: Stochastic Gradient Descent algorithm; regresses. Imports from miscFunctions.py, handles its own data.
- NeuralNetwork.py: Neural network; regresses and classifies. Imports from miscFunctions.py, activationFunctions.py, and dataHandler.py, has its data handled by dataHandler.py

- `LogisticRegression.py`: Logistic Regression algorithm; classifies. Imports from `activationFunctions.py` and `miscFunctions.py`, handles its own data.

Side scripts:

- `dataHandler`: Loads, generates, and handles datasets as well as handles plotting for `NeuralNetwork.py`,
- `miscFunctions.py`: Contains functions needed for functionality such as `to_categorical_numpy`, which transforms arrays into one-hot form, or `defaulting func`, which selects default values according to a priority list.
- `activationFunctions.py`: Contains activation functions, such as `softMax` or `ReLU_leaky` for use in logistic regression and the Neural Network.

Tested parameters are stored in the folder `parameterValues`. Analysis of the parameters is stored in text-form in `LaTeX/appendix/analyzedParameters` and in imageform in `LaTeX/images`. Source code files are stored directly in the main folder.

## 1.2 Franke function

The Franke function is the model used for linear regression. It is here defined for the domain  $x, y \in [0, 1]$

$$f(x, y) = \frac{3}{4} \exp \left( -\frac{(9x-2)^2}{4} - \frac{(9y-2)^2}{4} \right) + \frac{3}{4} \exp \left( -\frac{(9x+1)^2}{49} - \frac{(9y+1)^2}{10} \right) \\ + \frac{1}{2} \exp \left( -\frac{(9x-7)^2}{4} - \frac{(9y-3)^2}{4} \right) - \frac{1}{5} \exp \left( -(9x-4)^2 - (9y-7)^2 \right)$$

## 1.3 MNIST dataset

For classification methods, I use the MNIST image dataset. It contains 1797 entries of low-resolution images of the numbers 0-9 and accompanying labels describing what number the image portrays. there are 178 entries labelled 0 out of 1797 entries, making up 9.905% of the dataset, meaning any method that has a 9.905% accuracy is likely to have overflowed weights and is predicting every label to be 0.

## 1.4 Ordinary Least Squares (OLS) and Ridge

The OLS and Ridge methods, used in this project as comparison points for the linear regression, attempt to model a dataset by minimizing the square error between the dataset,  $z$ , and a model in the form  $\hat{z} = \mathbf{X}\theta$ , where  $\mathbf{X}$  is a design matrix and  $\theta$  is a vector of parameters.

OLS analytically minimizes the cost function  $C = \mathbb{E}[(z - \hat{z})^2]$  by using parameters

$$\theta = \left( \mathbf{X}^T \mathbf{X} \right)^{-1} \mathbf{X}^T y$$

Ridge amends OLS, attempting to mitigate issues brought by uninvertible or near-uninvertible design matrices, by adding a diagonal value before the inversion:

$$\beta = \left( \mathbf{X}^T \mathbf{X} + \lambda \mathbf{I} \right)^{-1} \mathbf{X}^T y$$

## 1.5 Stochastic Gradient Descent (SGD)

Stochastic Gradient Descent revolves around making prediction from a design matrix, initialized in some pre-determined way (here as random noise),  $X$ , to an expected value,  $z$ , with a 1D array,  $\theta$ , finding an error, and adjusting the array  $\theta$  with a gradient based on that error. The idea is that  $\theta$  will follow a gradient in the error-parameter space towards a location of an error minimum.

SGD as a method will attempt to regress and interpolate between datapoints to approach a continuous function.

The cost function used in this report for optimization is the MSE:  $C(\theta) = \frac{1}{n}(\mathbf{X}\theta - z)^T(\mathbf{X}\theta - z)$ . The gradient is hence given as  $\nabla C(\theta) = 2\mathbf{X}^T(\mathbf{X}\theta - z)$

### 1.5.1 Momentum

An option for optimizing the gradient descent is to implement a momentum feature. It functions as a running average with linearly increasing weights. Every iteration reduces the momentum by a factor and adds the current gradient to it;  $\text{momentum} \leftarrow \text{momentum} \times \rho + \nabla C(\theta)$ , where  $0 \leq \rho \leq 1$  is a drag factor

### 1.5.2 Minibatching

When optimizing on a gradient, there is no guarantee that there is only one minimum within the cost-space, and it is possible to be caught in local minima that have a larger error than the global minimum that we are searching for. To solve this, one could force more noise onto the gradient by adjusting for only one datapoint at a time, but this is very inefficient and removes parallelized computation as a possibility. In this project, minibatching is used to compromise between these. Minibatching calculates a gradient for a number of randomly selected points, making entrapment in local less likely while still maintaining efficiency. The cost function gradient can then be described as a sum of these points:  $\nabla C(\theta) = \sum_k \nabla_{\theta} c_k(\mathbf{x}_k, \theta)$ , where  $k$  is a subset of the original points referred to as the minibatch;  $k \in K \in N$ .

### 1.5.3 Varying step length

Step length is a factor multiplied onto the gradient that is applied to the weights in a descent-based method as a regularization tool. If the step-length is too small, the solution might be stuck in local minima or not find a solution in time. If the step length is too large, the solution might diverge. A varying step length can be implemented to move out of local minimum but still find accuracy within the minimum at the end of the training process. In this project, this is implemented as  $\eta = \frac{t_1}{t+t_0}$  where  $\eta$  is the step length,  $t_0$  and  $t_1$  are parameters to be adjusted, and  $t$  is an incrementing integer, starting at 0 and incrementing on each call.

## 1.6 Logistic Regression

Following a similar concept to SGD, logistic regression attempts to approximate a discrete dataset, typically being used as a classification tool. Where linear SGD will estimate the height of an entity, a logistic regressor will estimate whether that entity is a hill or a building.

In the instance of this project, the distinct classes are split into one-hot vector form and put through the Softmax function to determine what class has the best fit.

A logistic regression algorithm comes here in the form of (prediction =  $f(\mathbf{X}\theta)$ ), (error = prediction -  $z$ ), and ( $\theta \leftarrow \theta - X^T \text{error}$ ) where  $f()$  is the activation function, here Softmax.

### 1.6.1 $l_2$ regularization parameter

In order for weights not to diverge to overflow, regularization may be implemented in addition to the step length.  $l_2$  regularization refers to having an  $l_2$ -parameter between 0 and 1 that is multiplied by the sum of the squares of the weights to get an  $l_2$ -term, which is then added to the gradient. Since the gradient is subtracted from the weights, this can keep the weights from diverging.

## 1.7 Neural Network

An extension to logistic regression, a neural network places multiple layers between the input and the output, where each layer is considered an input to the following and output to the prior layer, allowing for modelling of much more complex problems.

## 2 Method

There are four methods tested, SGD, Ridge, Neural Network, and Logistic Regression. They have all been optimized by being given a parameter space, a combination of multiple parameters, having their accuracy tested on all of them, and selecting the best one. Further analysis is done on the parameters of this parameter-space.

### 2.1 SGD

To evaluate SGD, I implemented an SGD-algorithm with momentum, which was compared to analytical solutions Ordinary Least Squares (OLS) and Ridge. The algorithm was applied to a generated dataset made of a two-dimensional polynomial where both dimensions,  $x$  and  $y$ , where  $x, y \in [0, 1]$ . The polynomial was compared to the Franke function of those  $x$ - and  $y$ -values.

I then ran it through the all combinations of a set of parameters as given in the list below and found the optimal combination, and used those parameters to plot  $R^2$  and MSE as a function of the number of epochs.

The SGD data values were not split into training and testing sets, unlike the NN data, due to time constraints.

Parameters adjusted:

- Number of Datapoints in the dataset: 50, 100
- Polynomial degree of the dataset: 5, 8, 10, 12, 15
- Size of the minibatches: 10, 50, 100
- Number of epochs: 250, 500

- Different combinations of  $t_0$  and  $t_1$ : (10, 10), (10, 50), (10, 1)
- Momentum coefficient: 0.00, 0.25, 0.50, 0.75, 1.00

Finally, I looked at Ridge as a function of the  $\lambda$ -parameter to compare to my SGD solution.

## 2.2 Logistic Regression

To evaluate Logistic Regression, I implemented a logistic regression algorithm using momentum and an  $l_2$ -term. These values were evaluated against the accuracy of neural network that also was developed (subsection 2.3).

Similarly to the SGD-algorithm, the logistic regression-algorithm was tested through the parameter space, and the optimal set of parameters was selected from this.

The logistic regression data values were not split into training and testing sets, unlike the NN data, due to time constraints.

Parameters adjusted:

- Different combinations of  $t_0$  and  $t_1$ : (1, 1), (1, 3.16), (1, 10), (1, 31.6), (1, 100)
- Number of epochs: 50, 100, 500, 1000
- Size of the minibatches: 16, 32, 64, 128
- Momentum coefficients: 0.0, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0
- $l_2$  parameters: 10, 1, 0.1, 0.01, 0.001, 0.0

I then analyzed the accuracy score of the logistic regression as a function of the  $l_2$ -parameter and  $\eta$ .

## 2.3 Neural Network

The neural network was made before the logistic regressor, and the regressor was treated as a special case of a neural network, where there were 0 hidden layers. Since the neural network can both linearly regress and classify, it was compared to both SGD and logistic regression, spending 33 hours calculating the optimal parameters for the linear regression and 23 on the classification parameters. Since this is the main focus of the project, more time was spent analyzing the parameterspace for the NN than for the SGD-algorithm or the logistic regressor.

The biases were initialized at 0.01 and the weights were initialized as randomly distributed with a mean of 0 and a standard deviation of 1.

The softmax function was selected as output function for classification due to its normalization of the values given. For the linear regression, a linear function with a single parameter,  $\alpha$ , was used as a coefficient for the output. A linear function was chosen because a linear output can very easily approximate a linear function.

First I generated each combination of the parameters set and calculated an associated score as an average over 5 runs for each combination of the parameters. This was stored in a file. I then analyzed the score as a function

of the  $\lambda$ -parameter and the  $\eta$ -parameter, with  $\eta$  as both a constant and as a decreasing value (subsubsection 1.5.3). I then compared the different activation functions to each other.

Finally, I compared the performance of the neural network as a linear regressor to the performance of the SGD algorithm and Ridge and its performance as a classifier to the performance of the logistic regressor (subsection 3.4).

Parameters adjusted for linear regression:

- Number of hidden layers: 0, 1, 2, 4
- Number of nodes: 16, 32, 48, 64
- Number of epochs: 250, 500
- Size of minibatch: 32, 64
- $\eta$ , as  $t_0/(t + t_1)$ : (1, 1), (1, 10), (1, 100), (1, 1000), (1, 10000), (1, 100000), (1, 1000000)
- $\eta$ , as a constant: 1, 0.1, 0.01, 0.001, 0.0001, 0.00001, 0.000001
- $\lambda$ : 0.1, 0.01, 0.001, 0.0001, 0.00001, 0.000001
- $\alpha$ : 0.1, 0.01, 0.001
- Activation function for the hidden layers: sigmoid, ReLU\_leaky, ReLU
- Activation function for the output layer: softmax

Parameters adjusted of classification:

- Number of hidden layers: 0, 1, 2
- Number of nodes: 16, 32
- Number of epochs: 250, 500, 750
- Size of minibatch: 32, 64, 128
- $\eta$ , as  $t_0/(t + t_1)$ : (1, 1), (1, 10), (1, 100), (1, 1000), (1, 10000), (1, 100000), (1, 1000000)
- $\eta$ , as a constant: 1, 0.1, 0.01, 0.001, 0.0001, 0.00001, 0.000001
- $\lambda$ : 0.1, , 0.00316, 0.000001
- $\alpha$ : 1, 0.1
- Activation function for the hidden layers: sigmoid, ReLU\_leaky, ReLU
- Activation function for the output layer: softmax

For the neural network, the data is split into training and test data (75/25 split), where the data is trained on the training data and tested on the test data. This prevents overfitted functions from receiving too high a score.

### 3 Results/Discussion

#### 3.1 SGD

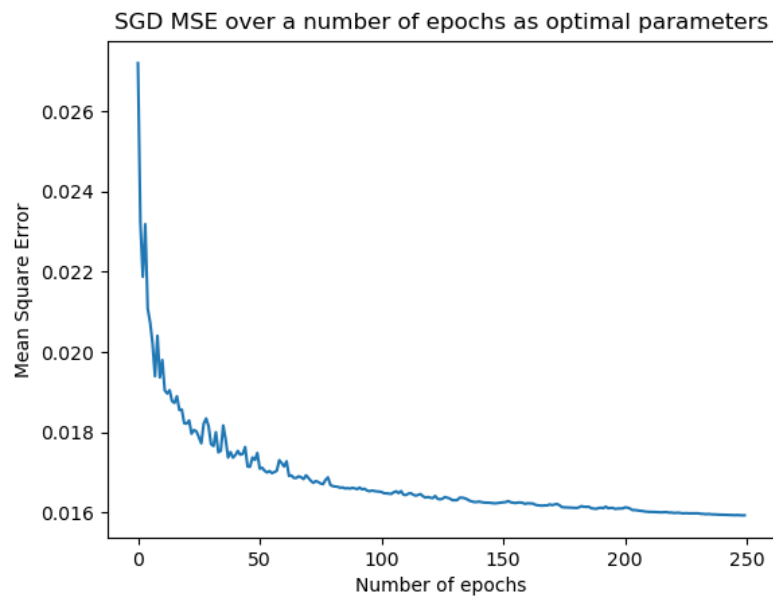


Figure 1: Shows how well the SGD solution approaches the Frankefunction as a function of the number of epochs by using the MSE of the training data.

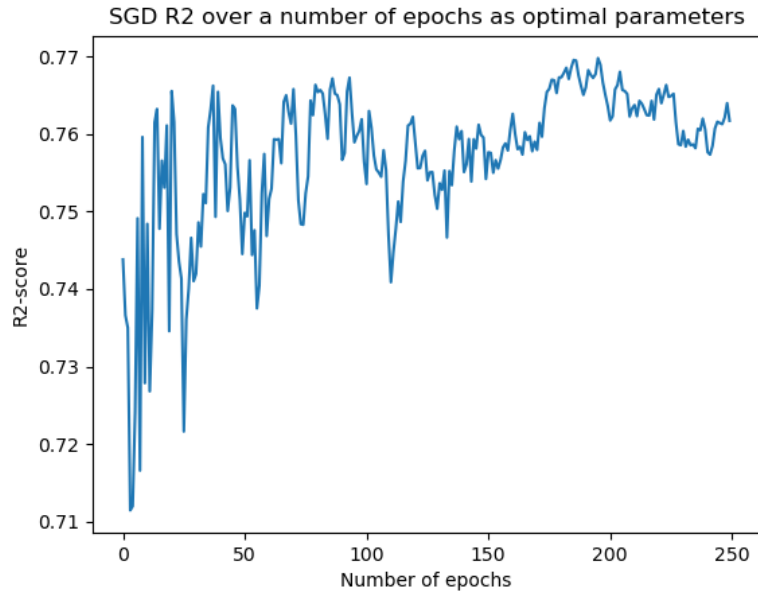


Figure 2: Shows how well the SGD solution approaches the Frankefunction as a function of the number of epochs by using the  $R^2$ -score of the training data.




Figure 3: Shows how well the ridge solution approaches the Frankefunction as a function of the  $\lambda$ -parameter by using the  $R^2$ -score.



Figure 1 and Figure 2 show that, perhaps unsurprisingly, a greater number of epochs leads to a higher accuracy, although the  $R^2$  seems to vary a large amount. This is perhaps due to the cost function being based on the MSE and not the  $R^2$ .

The fact that this is all training data and not test data suggests that the MSE and  $R^2$  should not be trusted too highly at a higher number of epochs.

Table 1: A sample of  $\lambda$ -values with an associated  $R^2$  score.

$\lambda$	$R^2$
1.0	0.922590
0.117681	0.952397
0.0001	0.974680
0.0/OLS	0.979087

The analytical Ridge-solution reaches much higher  $R^2$ -scores than the numeric SGD, suggesting that a simple numeric approach is not good enough to model the Franke function. The analytical solution performs best as a standard OLS, or with  $\lambda = 0$ .

Table 2: Shows score reached as a function of  $\eta$ . An  $\eta$ -value containing a "None" means that the other value is considered as a constant eta value.

Eta-values	Average $R^2$	Median $R^2$	Maximum $R^2$	Non-overflown SGD's [%]
('10', '10')	0.688806	0.742613	0.797290	1.0
('10', '1')	0.620894	0.759868	0.793156	1.0
('10', '50')	0.680007	0.737220	0.799657	1.0
('100', '1')	-0.007316	-0.001517	-6.968662e-08	0.9375
('50', '1')	-0.005946	-0.000307	-4.494689e-10	0.979167
('None', '0.1')	0.820982	0.874075	0.925629	0.943333
('None', '0.01')	0.786223	0.736666	0.884543	1.0
('None', '0.001')	0.646276	0.810326	0.844811	1.0

Table 3: Optimal values found with associated  $r^2$  score. Franke was generated with 200 points on each axis, and a design matrix of polynomial degree 15 was used.

$R^2$	(t0, t1)	# of epochs	Minibatch size	Momentum-coefficient
0.760531	(10, 1)	1000	128	0.0

Table 2 seems to strongly favour constant  $\eta$ -values, as is also reflected in subsection 3.3. This is unexpected and may be due to poorly chosen  $t_0$  and  $t_1$  values, but due to constant values performing better for every value, this seems unlikely.

## 3.2 Logistic Regression

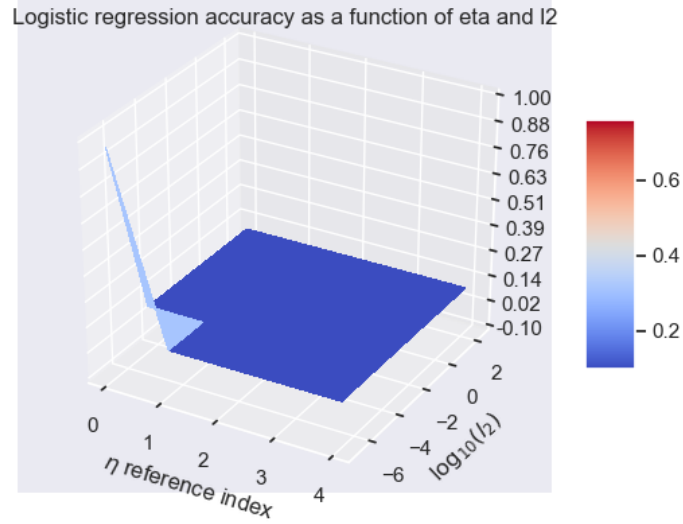


Figure 4: shows how well the logistic regression predicts the correct labels as a function of the  $l_2$ - and  $\eta$ -parameter. The  $\eta$ -reference index refers to  $t_0$  and  $t_1$  in order as (1.0, 100.0), (1.0, 31.622777), (1.0, 10.0), (1.0, 3.162278), (1.0, 1.0).

Table 4: Optimal values found with associated accuracy score.

Accuracy	( $t_0$ , $t_1$ )	# of epochs	Minibatch size	Momentum-coefficient	$l_2$
0.987201	(1, 31.622777)	1000	128	0.9	0.0

Figure 4 reflects a poor choice in  $\eta$  and  $l_2$  values, as most of the graph shows an accuracy of 9.905% which implies, as mentioned in subsection 1.3, that the weights overflowed. Table 2 shows an error rate of 1.28%, which is a fairly high performance when compared to the values found at [1]

## 3.3 Neural Network

Table 5:  $R^2$  score reached for each activation function

Function	Average	Median	Maximum	Non-overflown NN's [%]
ReLU_leaky	-37160.106071	-1.620086	0.900225	0.982143
ReLU	-43070.479707	-1.975616	0.863374	0.982143
sigmoid	-5.518974+194	-1.202853	0.882485	0.943204

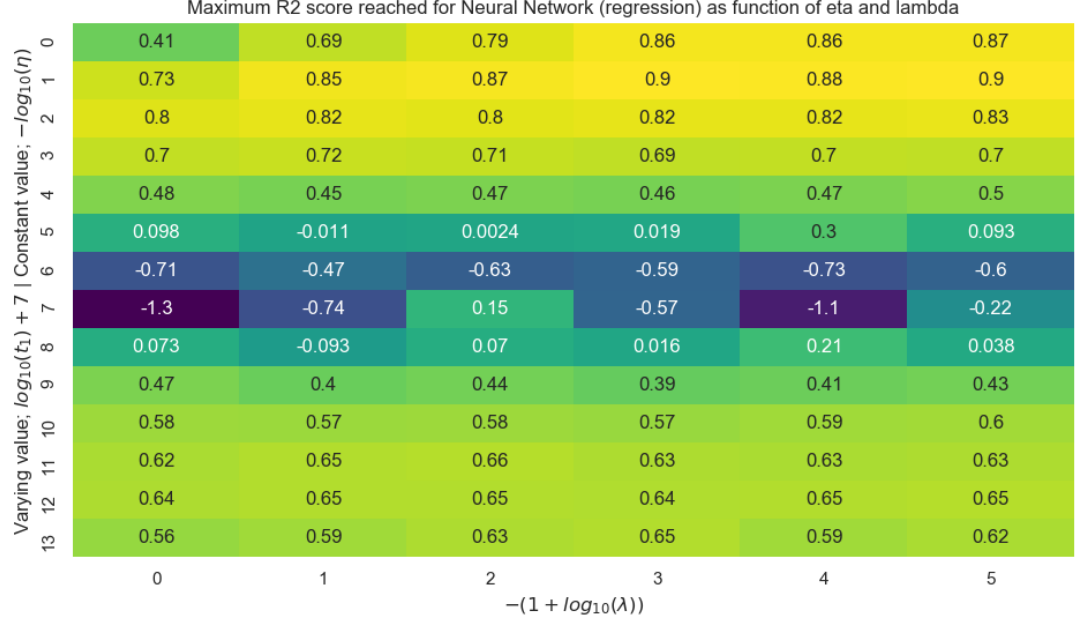


Figure 5:  $R^2$  score as a function of  $\eta$  and  $\lambda$ . The axes are numbered by indexes, not actual axis value, but axis labels describe the conversion between the values and the indexes. The  $\eta$ -axis is split into constant  $\eta$ , indices 0 to 6 inclusive, and varying  $\eta$ , indices 7 to 13 inclusive. Varying  $\eta$  assumes  $\eta = t_0/(t + t_1)$  where  $t_0 = 1$

From Table 5, ReLU\_leaky manages to outperform both ReLU and sigmoid in nearly every aspect (scoring equally well to ReLU in number of overflowed NN's), and was found to be the optimal activation function for both classification and linear regression. Meanwhile, whilst sigmoid outperformed ReLU at its best and at its median, it overflowed more networks, and thus seems less stable or reliable.

From Figure 5, a constant  $\eta$  performs better than a varying  $\eta = t_1/(t + t_0)$ , optimally at  $\eta = 0.1$ , paired with  $\lambda = 1e - 4$  or  $1e - 6$ .

Table 6: Optimal parameters with an increased number of epochs. The classification network was trained for 1000 epochs with a minibatch of 32 datapoints, while the linear regression network was trained for 5000 epochs with a minibatch of 32 datapoints.

Mode	Score	Layers	Nodes	$(t_0, t_1)$	$\lambda$	$\alpha$	activFunc	outFunc
Classification	0.971	0	32	(1, 1e3)	0.1	1.0	ReLU_leaky	softmax
Linear regression	0.994	1	32	0.1	1e-4	0.1	ReLU_leaky	linearActivation

## 3.4 Comparing the algorithms

### 3.4.1 Regression

The neural network manages to attain an  $R^2$  score of 0.994 (Table 6) compared to SGD's 0.760 (Table 3) and OLS' 0.979 (Table 1). This makes a neural network a much stronger tool for linear regression than basic SGD for the Franke function, and outperforms analytical Ridge and OLS. This should be contrasted with the difficulty of implementation and optimization, where OLS is very easy to implement and takes no resources to optimize (just training), compared to the slightly harder SGD and much harder NN.

### 3.4.2 Classification

The neural network manages to attain an accuracy score of 0.971 (Table 6), compared to the logistic regressions 0.987 (Table 4), indicating that a simple logistic regression outperforms a neural network, even when that network has 0 layers, emulating a simple logistic regression algorithm. Since the optimal logistic regression  $l_2$ -parameter was found to be 0, this is likely to be due to the momentum feature implemented into the logistic regression, having a momentum-coefficient of 0.9. I expect that if the  $l_2$ -parameter and the momentum-feature were to be implemented into the neural network, it would not be outperformed by the logistic regression algorithm, as it should be able to emulate its performance with 0 hidden layers.

## 3.5 Improvements

In this project I have exclusively used the MNIST dataset and the Franke function, and I would be able to get a stronger basis for any conclusions if I spread the analysis over more datasets or functions.

## 4 Conclusion

The neural network outperforms all other tried methods (SGD, Ridge, and OLS) for linear regression based on  $R^2$ -score. The neural network is outperformed by a logistic regression-algorithm when used for classification, but this is assumed to be due to additional features being implemented into the logistic regression (momentum and  $l_2$ -regularization), as a neural network should be able to emulate a logistic regression by using 0 hidden layers. The expectation is that if those features were to be implemented into the neural network, it would be able to outperform all other methods tested.

## 5 References

- 1 Y. LeCun, L. Bottou, Y. Bengio and P. Haffner: *Gradient-Based Learning Applied to Document Recognition*, Proceedings of the IEEE, 86(11):2278-2324, November 1998, retrieved from <http://yann.lecun.com/exdb/mnist/>

## 6 Appendix